FPL AI:
A REINFORCEMENT LEARNING AGENT FOR FANTASY FOOTBALL

Tony Munene Kinyua
munene.kinyua@yale.edu

Advisor: James Glenn
james.glenn@yale.edu

A Senior Thesis as a partial fulfillment of requirements for the Bachelor of Science in Computer Science and Economics

Department of Computer Science
&
Department of Economics
Yale University
May 1, 2025

Acknowledgements

Contents

A REINFORCEMENT LEARNING AGENT FOR FANTASY FOOTBALL

Tony Munene Kinyua

Abstract

The abstract is a brief description of the paper. Write it in 100-150 words in general. For the thesis, it can be longer. Present your research questions, your contribution to the computer science literature, a concise description of your work, and how you answered your research question. Add a statement of the main results. This part can be revisited when the paper is complete.

Then, you can revisit the title of the paper. If you are revisiting the title of your paper, ask this question: Which one best explains the main point of the paper?

- Microfinance

- Microfinance and Economic Development

- Microfinance as a Strategy for Economic Development: Realizing its Potential for Improving the Standing of Women

Now, examine your title. Does it show the main point of the paper? If you wanted to rewrite it, how would you do it? Review Lecture 5: How to Research (II), Senior Thesis Structure for details in writing your thesis.

1 Introduction

The task of picking a team for Fantasy Premier League (FPL) is an ardous one to say the least. One has to balance their bias for the team they support and their analysis of high performing players. Furthermore, the time difference between the United States and England makes it difficult for supporters of English Premier League (EPL) teams to watch live matches since they're usually played very early on the weekends. This becomes even more difficult for students and the working class who dedicate Saturday mornings to sleeping in after a long week of work. As such, many a FPL manager drop out from active team management because they don't have enough time to keep up with the matches, transfer news, injury updates, gameweek restructuring news, and other news that require a religious following of the EPL. This is why I came up with this passion project of building an agent that could actively manage an FPL team so that you can spare your weekend mornings to catching up with sleep.

How well can an agent, trained through reinforcement learning, perform compared to the average FPL manager? FPL is a social sport. Where's the fun in playing FPL if you don't play against your mates? Be it in a classic league or a head-to-head one, I always look forward to seeing where I place at the end of a gameweek. Therefore, I just can't let any agent run my FPL team. There's too much on the line. The agent was trained on data from 2016/17 season to 2022/23 and evaluation was done using data from the last complete season - 2023/24.

Essentially, the introduction is a summary of the paper. A great introduction is one that makes the reader excited to read the rest of the paper. A good introduction encourages readers to read your work with interest and prepares them to understand it better.

Remember the three-pass reading that we learned in Lecture 2 of the class, CPSC 490 Senior Project? When you write a paper, you can expect most reviewers (and readers) to make only one pass over it. (i.e. 5 min, max 10 min.) Then, decide whether they will read more or put it aside.

A strong introduction does the following in this general order:

- Motivates why the subject of the research is important

- States the research question

- Discusses where the paper fits in existing literature

- Describes the contribution of the paper

- States the research methods

- States the main results

In this template, section titles are provided as suggestions, and you are more than welcome to change them, except for Abstract, Introduction, Background, Related Work, and Conclusion, which are must-have.

Please share the document with your advisor with edit previllage. Comments can be provided *sk: this way, and comments can be removed when addressed.

2 Background

2.1 Background: English Premier League and Fantasy Premier League

2.1.1 The English Premier League: Structure and Significance

The English Premier League (EPL) is the top tier of professional football (soccer) in England, founded in 1992 after breaking away from the Football League [1]. It consists of 20 clubs that compete in a double round-robin tournament, playing 38 matches each season (home and away against every other team). The season typically runs from August to May, with teams awarded three points for a win, one for a draw, and none for a loss. At the end of each season, the three lowest-ranked teams are relegated to the Championship (second tier), while three teams are promoted from the Championship to the Premier League [2].

The EPL has grown to become the most-watched sports league globally, broadcasting to 212 territories with a potential audience of 4.7 billion people [3]. Its commercial success is unprecedented, with the 2022-2025 broadcasting rights valued at approximately £10 billion [4]. This financial power has enabled EPL clubs to attract elite players and coaches from around the world, contributing to the league's competitive nature and global appeal [5].

2.1.2 Fantasy Premier League: Game Mechanics and Popularity

Fantasy Premier League (FPL) is the official fantasy sports game associated with the English Premier League. Launched in 2002, it has grown exponentially to over 11 million players worldwide as of the 2023/24 season [6]. FPL allows participants to assemble a virtual team of real Premier League players within specific constraints and earn points based on those players' actual performances in Premier League matches [7].

Basic Rules and Structure

Participants (known as "managers") are allocated a virtual budget (£100 million) to select a 15-player squad consisting of:

- 2 Goalkeepers

- 5 Defenders

- 5 Midfielders

- 3 Forwards

The budget constraint forces managers to balance premium-priced elite players with cheaper options. Each gameweek, managers select 11 players from their 15-player squad to form a starting lineup, with the remaining 4 players on the bench. Additional constraints include:

- Maximum of 3 players from any single Premier League club

- Formation requirements (minimum of 1 goalkeeper, 3 defenders, and 1 forward)

- Limited free transfers between gameweeks (typically 1 per week, with additional transfers costing points) [8]

Scoring System

Points are awarded based on players' real-world performance metrics:

- Appearance (playing at least 60 minutes): 2 points

- Goals: 6 points (midfielder), 4 points (forward), 6 points (defender/goalkeeper)

- Assists: 3 points

- Clean sheets: 4 points (defender/goalkeeper), 1 point (midfielder)

- Saves: 1 point per 3 saves (goalkeeper)

- Penalties saved: 5 points (goalkeeper)

- Bonus points: 1-3 additional points to the top performers in each match

Negative points are also assigned for:

- Yellow cards: -1 point

- Red cards: -3 points

- Own goals: -2 points

- Penalties missed: -2 points

- Goals conceded: -1 point per 2 goals (defender/goalkeeper) [8]

Special Features

FPL includes several strategic elements that increase its complexity:

- **Captain**: Managers designate one player as captain each gameweek, doubling their points

- **Vice-captain**: A backup who becomes captain if the original captain doesn't play

- **Chips**: Special boosts used once per season:

    - Bench Boost: Points from bench players count for one gameweek
    - Triple Captain: Triple (rather than double) points for the captain
    - Free Hit: Unlimited free transfers for one gameweek only
    - Assistant Manager: Add a manager to your team to score points for three consecutive gameweeks
    - Wildcard: Unlimited free transfers that permanently change the team [8]

2.1.3 Data and Performance Metrics in Football

Traditional Statistics

Football has historically relied on basic statistics to evaluate performance:

- Goals and assists

- Clean sheets

- Shots and shots on target

- Pass completion percentage

- Possession percentage

- Cards and fouls [9]

Advanced Metrics

Recent years have seen an explosion in advanced metrics:

- Expected Goals (xG): Probability of a shot resulting in a goal

- Expected Assists (xA): Probability of a pass leading to a goal

- Progressive Passes/Carries: Passes/carries that move the ball significantly toward the opponent's goal

- Defensive Actions: Tackles, interceptions, clearances, and blocks

- Pressure Events: Instances of applying pressure to an opponent

- VAEP (Value of Actions by Estimating Probabilities): Calculating the value of every action [10, 11]

Player Pricing and Value

FPL assigns each player a monetary value, which fluctuates throughout the season based on ownership patterns. The game adjusts player prices according to transfer market dynamics:

- Players transferred in by many managers typically increase in price

- Players transferred out by many managers typically decrease in price

- Price changes occur in £0.1m increments within certain thresholds [12]

This dynamic pricing creates a parallel "market economy" that influences decision-making, as managers must consider not only point-scoring potential but also value appreciation/depreciation [13].

2.1.4 Decision-Making Challenges in FPL

Team Selection Complexity

The fundamental challenge in FPL is optimizing team selection under constraints. With approximately 500 Premier League players available, the theoretical number of valid 15-player squads exceeds 10^{23}. Even limiting to weekly starting 11 selections, the decision space remains enormous [14].

Predictive Uncertainty

Football is inherently unpredictable, with significant variance in player performance. Key uncertainties include:

- Injuries and rotation (players rested for certain matches)

- Form fluctuations throughout the season

- Managerial decisions affecting player roles and playing time

- Match context and fixture difficulty

- Weather conditions and other external factors [15, 16]

Multi-objective Optimization

FPL managers must balance competing objectives:

- Maximizing expected points for the current gameweek

- Planning for future gameweeks (favorable fixture runs)

- Building team value through strategic transfers

- Differential selection (picking low-ownership players for competitive advantage)

- Risk management (captaincy choices, bench quality) [17]

Temporal Dynamics

The game spans 38 gameweeks, requiring both short and long-term planning:

- Weekly decisions: Starting lineup, captaincy, transfers

- Medium-term decisions: Chip usage, planning for blank/double gameweeks

- Season-long decisions: Overall strategy and style of play [18]

2.1.5 Relationship to Reinforcement Learning

Fantasy Premier League presents an ideal environment for reinforcement learning applications due to several characteristics:

Markov Decision Process Formulation

FPL naturally fits into the Markov Decision Process framework:

- **States**: Current team composition, budget, available transfers, fixture schedule

- **Actions**: Transfers, captain selection, bench order, chip usage

- **Transitions**: How actions transform the state (affected by real-world player performances)

- **Rewards**: Gameweek points earned

- **Long-term rewards**: Season-long point accumulation [19, 20]

Delayed Rewards and Credit Assignment

FPL exhibits the classic reinforcement learning challenge of delayed rewards:

- Transfer decisions may not pay off immediately

- Building team value early may enable stronger teams later

- Planning for fixture difficulty must account for weeks or months ahead [21]

Exploration-Exploitation Tradeoff

Successful FPL strategy requires balancing:

- Exploitation: Selecting proven performers and popular captaincy options

- Exploration: Taking calculated risks on differentials or emerging players [22]

Non-stationarity

The FPL environment is non-stationary due to:

- Player form changes throughout the season

- Team tactical evolutions

- Injury impacts

- Transfer windows (January) bringing new players

- Manager changes affecting team performance [23]

2.1.6 Previous Research and Algorithmic Approaches

Optimization-Based Approaches

Early algorithmic approaches to FPL focused on optimization techniques:

- Linear programming for team selection

- Integer programming for transfer planning

- Mixed-integer programming for season-long planning

While effective for constrained selection problems, these approaches often struggle with the inherent uncertainty and temporal dynamics of football [24, 25].

Machine Learning Applications

Recent research has increasingly applied machine learning:

- Regression models for player point prediction

- Time series forecasting for form prediction

- Classification models for clean sheet probability

- Ensemble methods combining multiple prediction approaches [26, 27]

Reinforcement Learning Explorations

Emerging research applies reinforcement learning to FPL:

- Q-learning for transfer decisions

- Deep Q-Networks for team selection

- Policy gradient methods for season-long strategy

- Monte Carlo Tree Search for planning [28, 29]

These approaches show promise in managing the complex, sequential decision-making process that FPL represents, while accounting for uncertainty and delayed rewards.

2.1.7 Data Sources and Availability

Modern FPL research benefits from unprecedented data availability:

- Official FPL API providing comprehensive game data

- Third-party websites aggregating historical performance

- Event-level data from commercial providers (Opta, StatsBomb)

- Community resources like public GitHub repositories of historical data

- Web scrapers that collect and organize player statistics [30, 31]

This rich data ecosystem enables the training of sophisticated models that can make informed predictions about player performance and optimal decision strategies.

3 Methodology

3.1 Introduction to Methodology

My methodological apprach in this project can be divided into three: modeling beliefs on players' abilities in a Bayesian manner, modeling FPL's team selection process as a Belief-State Markov Decision Process (BSMDP), and solving the MDP using Bayesian Q-learning. This methodology is similar to the one described in the paper - Competing with Humans at Fantasy Football: Team Formation in Large Partially-Observable Domains [14]. I have, however, made my own adjustments to this methodology in ways that I believe improved my model.

3.2 Modeling players' abilities

Modeling a player's ability in a manner that captured the uncertainity of their performance in games was important in this project. As such, I took a Bayesian apprach to model this belief by maintaining a distribution over possible performance levels. Furthermore, a Bayesian model allowed me to incorporate domain knowledge through priors using performance data from previous seasons. This was crucial in giving good players a higher baseline even when they experience a slump in their performance. The players' abilities I sought to model are the probabilities of a player:

- scoring a goal

- assisting a goal

- starting a game

- getting subbed during a game

- remaining unused during a game

Defining the following terms is essential in describing my methodology. For the i-th gameweek, we define:

- M_i as the set of matches in gameweek i.

- P_i as the set of players available for selection in gameweek i.

- A_i as the set of actions available in gameweek i, where $a \in A_i$ is a subset of P_i and observes all team selection constraints.

- $p_i \in P_i$ is associated with its FPL-designated position $pos(p_i)$ and price $pr(p_i)$.

- $\tau_p \in \tau$ is a system of distributions representing the player's performance/influence on the matchplay.

- O_i is the set of match observations in gameweek i.

- $o \in O_i$ includes both the result of the matches and the performance of the players in the selected team e.g. goals, assists, clean sheets, yellow cards, red cards, bonus points. The probability of each $o \in O_i$ is somehow dependent on the players' characteristics (τ) i.e. a team with strong attackers is more likely to score goals, therefore, $P(o|\tau)$ is dependent on τ.

- $R(o, a_{prev}, a_{curr})$ is the reward function, which returns the points scored by the selected team a_{curr}, given the match observations o. The previous team a_{prev} is also provided to penalize the agent for any poor player transfers or transfers beyond the allowed number.

I use three distributions to model players' abilities:

- ρ_p - a three-state categorical distribution representing the player's probability of starting a match, being substituted, or not playing at all i.e. (start, sub, unused).

- ω_p - a Bernoulli/Binomial distribution over a single trial, representing the probability of a player scoring a goal given he was playiong at the time

- ψ_p - a Bernoulli distribution representing the probability of a player providing an assist given he was playing at the time

Using a Bayesian approach allowed me to leverage the respective distributions' conjugates to update the players' priors (belief) using data from previous seasons. I defined uniform priors for all players as described in [14] as follows:

$$\omega_p \sim Beta(1,1), \psi_p \sim Beta(1,1), \rho_p \sim Dirichlet(\frac{1}{4}, \frac{1}{4}, \frac{1}{4})$$

I further defined four multinomial distributions S_{pos}, one for each position - to describe the how long players who play the same position are likely to play, given they start a match. These distributions were defined using a Dirichlet distribution, modeling the probability a player from the respective position pos leaving the match at minute x, where $0 \leq x \leq 90$.

Samples of a player's ability, τ_p, and minutes played in a game, given they were in the starting lineup, is drawn from these conjugate distributions

I simulate a gameweek by simulating each fixture in the gameweek as follows (The procedure focuses on the home team for conciseness but is also applicable to the away team):

Team	Formation
AVL, BHA, BOU, CHE, FUL, MCI, MUN, TOT, WHU	4-2-3-1
ARS, CRY, LIV, NEW	4-3-3
LUT, WOL	3-4-2-1
BRE, SHU	3-5-1
EVE	4-4-1-1
BUR (classic)	4-4-2
NFO	4-2-3-1

Table 3.1: Favored formations for the 2023/24 Premier League Season

- Define P_H and P_A as the set of players available the home and away teams respectively. I used formation frequency data from the English Premier League to determine individual team compositions. As such, I assigned each team a default formation as follows:

- I, however, had to make some modifications when constituting these teams using the aforementioned formations. In the case where a team does not have enough forwards to fill the required number as per their assigned formation (as was the case with Totttenham and Newcastle), I used midfielders instead. Further, since Premier League data does not distinguish between attacking and defensive midfielders, I simplified formations with such distinctions e.g. 3-4-2-1 and 4-2-3-1 by grouping both as general midfielders.

- Sample τ_p for each player $p \in P_H$ from the belief model $Pr(\tau_p|b_i)$

- Randomly select eleven players from P_H in proportion to their probability of starting the match i.e. $Pr(\rho_p = start)$ These players constitute the starting lineup L_H

- The minute each player leaves the pitch is sampled from the S_{pos} distribution for the player's position pos

- Each player in P_H and not in L_H is assigned to the set of substitutes U_H

- For every minute that a player in L_H is set to get substituted:

- We randomly select a player from U_H to replace the outgoing player in proportion to the probability of the player being substituted i.e. $Pr(\rho_p = sub)$

- The replacement is added to L_H (removed from U_H). We further assume that the player being substituted is not substituted again in the same match.

- We use the Dixon-Coles model [23] predict the outcome of the fixture. The model extends the basic Poisson model for soccer prediction by assuming that goals scored by teams follow a Poisson distribution. It also accounts for team-specific attacking and defensive threats, and home advantage while adding a crucial correction for the

dependency between team's scores, especially for low-scoring results (0-0, 1-0, 0-1, 1-1) I was fortunate to find a clean implementation of the Dixon-Coles model on David Sheehan's article on Predicting Football Results with Statistical Modelling: Dixon-Coles and Time-Weighting* [32]

- If a goal is scored, it is allocated to player p with probability $Pr(\omega_p = 1)$ while an assist is allocated to player p with probability $Pr(\psi_p = 1)$. I assume that every goal has attributed assist.

- Other point scoring guidelines i.e. scoring minutes played and clean sheets proceed at described in 2.1.2

These point estimates were used in combination with the BSMDP reward function R to approximate the immediate reward from performing an action

3.3 Modeling the FPl team selection problem

Similar to the prior problem of modeling players' abilities, selecting an FPL team faces the fundamental problem of making decisions under uncertainity. One doesn't know whether a player will start a game, get injured, or even how long they will play. This is why I opted for a Belief-State Markov Decision Process (BSMDP) as opposed to the standard MDP. The former assumes perfect knowledge of states.

Fpl is inherently sequential - decisions in gameweek 1 affect options in gameweek 2 and beyond due to budget constraints, free transfer limitations, and team value changes based on player price fluctuations. As such, the BSMDP naturally captures the sequential nature of FPL team selection. I also reinfor

The Belief State Markov Decision Process is defined by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \gamma)$, where:

- $\mathcal{S}$ is the state space

- $\mathcal{A}$ is the action space

- $\mathcal{T} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ is the transition function

- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ is the reward function

- $\mathcal{O}$ is the observation space

- $\gamma \in [0, 1)$ is the discount factor

In a belief state MDP, the agent maintains a belief distribution over possible states rather than knowing the exact state. We will now formalize each component in the context of the FPL environment.

3.3.1 State Space

The state space $\mathcal{S}$ in the FPL environment is multi-dimensional and consists of:

$$\mathcal{S} = \{(B, \mathbf{T}, GW, \mathbf{P})\} \tag{3.1}$$

Where:

- $B \in \mathbb{R}^+$ represents the remaining budget (with initial value $B_0 = 100.0$)

- $\mathbf{T} \in \mathcal{P}_{15}$ represents the current team of 15 players, where $\mathcal{P}$ is the set of all available players

- $GW \in \{1, 2, ..., 38\}$ represents the current gameweek

- $\mathbf{P} \in \mathbb{R}^{n \times 2}$ represents performance predictions for all n players

Each player $p \in \mathcal{P}$ has attributes including:

- Position $pos(p) \in \{\text{GK}, \text{DEF}, \text{MID}, \text{FWD}\}$

- Team $team(p) \in \{1, 2, ..., 20\}$

- Price $price(p) \in \mathbb{R}^+$

- Expected points $E[points(p, gw)] \in \mathbb{R}^+$ for each gameweek gw

3.3.2 Action Space

The action space $\mathcal{A}$ consists of three main components:

$$\mathcal{A} = \mathcal{A}_{transfer} \times \mathcal{A}_{captain} \times \mathcal{A}_{bench} \tag{3.2}$$

Where:

- $\mathcal{A}_{transfer} \subset \mathcal{P} \times \mathcal{P}$ represents possible player transfers (sell, buy)

- $\mathcal{A}_{captain} \subset \mathcal{P}$ represents the choice of captain from the team

- $\mathcal{A}_{bench} \subset \binom{\mathcal{P}}{4}$ represents the choice of 4 players to bench

In the implementation, the action space is simplified to selecting among a subset of promising actions as suggested in [14]. In each iteration of training the reinforcement learning agent, we replaced the weakest member of the simplified action set with a promising member of the unexplored action space. I defined a player p's $value$ as their scored points per unit cost i.e. $\frac{points(p)}{price(p)}$

$$\mathcal{A}_{simplified} = \{0, 1, 2\} \tag{3.3}$$

where each index corresponds to a dynamically maintained action subset.

3.3.3 Constraints

The FPL environment imposes several constraints as described in 2.1.2

3.3.4 Transition Dynamics

The transition function $\mathcal{T}$ for the FPL environment can be decomposed as follows:

$$\mathcal{T}((B, \mathbf{T}, GW, \mathbf{P}), a, (B', \mathbf{T}', GW', \mathbf{P}')) = \begin{cases} 1 & \text{if conditions are met} \\ 0 & \text{otherwise} \end{cases} \tag{3.4}$$

Where the conditions are:

$$GW' = GW + 1 \tag{3.5}$$
$$\mathbf{T}' = (\mathbf{T} \setminus \{p_{sell}\}) \cup \{p_{buy}\} \text{ if } a \text{ includes a transfer} \tag{3.6}$$
$$B' = B + price(p_{sell}) - price(p_{buy}) \text{ if } a \text{ includes a transfer} \tag{3.7}$$
$$\mathbf{P}' = f(GW') \text{ (updated player performance predictions)} \tag{3.8}$$

The transition is deterministic given the action and the player performance predictions.

3.3.5 Reward Function

The reward function $\mathcal{R}$ is defined as the points earned in a gameweek minus any transfer penalties. The captain's points are counted twice as per FPL special features 2.1.2

$$\mathcal{R}((B, \mathbf{T}, GW, \mathbf{P}), a) = \sum_{p \in \mathbf{T}_{playing}} points(p, GW)$$
$$+ points(captain, GW) - transfer\_cost \tag{3.9}$$

Where:

- $\mathbf{T}_{playing} \subset \mathbf{T}$ is the subset of 11 players not on the bench

- $captain \in \mathbf{T}_{playing}$ is the selected captain

- $transfer\_cost = 4 \times \max(0, num\_transfers - free\_transfers)$

3.4 Solving the FPL team selection Problem

Bayesian Q-learning is particularly well-suited for solving the FPL BSMDP since it directly incorporates uncertainty about the value function itself. Rather than maintaining a point estimate of Q-values as in standard Q-learning, it maintains a probability distribution over possible Q-values. This helps deal with uncertainities about the true value of

performing a transfer action $a \in \mathcal{A}$. Further, it makes better use of the relatively few data points (38 gameweeks), by incorporating prior knowledge from previous gameweeks.

For each potential action $a = (p_{sell}, p_{buy})$, the agent maintains a belief distribution over the Q-value as follows:

$$Q(s, a) \sim \mathcal{NG}(\mu_a, \lambda_a, \alpha_a, \beta_a) \tag{3.10}$$

Where $\mathcal{NG}$ is a Normal-Gamma distribution with:

- μ_a: mean estimate of the Q-value

- λ_a: precision parameter

- α_a: shape parameter

- β_a: rate parameter

3.4.1 Bayesian Q-Value Update

After taking action a and observing reward r, the belief distribution is updated according to the normal-gamma update rules:

$$\lambda'_a = \lambda_a + 1 \tag{3.11}$$
$$\alpha'_a = \alpha_a + 0.5 \tag{3.12}$$
$$\mu'_a = \frac{\lambda_a \mu_a + r}{\lambda'_a} \tag{3.13}$$
$$\beta'_a = \beta_a + \frac{0.5\lambda_a(r - \mu_a)^2}{\lambda'_a} \tag{3.14}$$

3.4.2 Value of Perfect Information (VPI)

The environment uses the Value of Perfect Information (VPI) to balance exploration and exploitation as suggested in [14]. For each action a, the VPI is calculated as:

$$VPI(a) = \begin{cases} \sigma_a \cdot t_\nu(z) \cdot (1 - CDF_\nu(z)) + \sigma_a \cdot PDF_\nu(z) & \text{if } a = a^* \\ \sigma_a \cdot z \cdot CDF_\nu(z) + \sigma_a \cdot PDF_\nu(z) & \text{if } a \neq a^* \end{cases} \tag{3.15}$$

Where:

- a^* is the action with the highest estimated mean Q-value

- $\nu = 2\alpha_a$ is the degrees of freedom for the t-distribution

- $\sigma_a = \sqrt{\frac{\beta_a(1 + 1/\lambda_a)}{\alpha_a}}$ is the standard deviation

- $z = \frac{Q(a') - \mu_a}{\sigma_a}$ where $Q(a')$ is the Q-value of the best alternative action if $a = a^*$, or the Q-value of the best action if $a \neq a^*$

- CDF_ν and PDF_ν are the cumulative distribution function and probability density function of the t-distribution with ν degrees of freedom

3.4.3 Action Selection and Exploration

The action selection mechanism combines exploitation (choosing the action with the highest estimated Q-value) with directed exploration using VPI:

$$a_{selected} = \arg\max_a \{\mu_a + VPI(a)\} \tag{3.16}$$

Additionally, the environment dynamically updates the action subset by replacing actions with low utility (defined as $\mu_a + VPI(a) < \mu_{a^*}$) with newly generated promising actions as described in 3.3.2

3.4.4 Algorithm

The initial team is selected greedily based on expected points per unit cost:

$$value(p) = \frac{E[points(p, 1)]}{price(p)} \tag{3.17}$$

The team is selected to maximize the sum of values while respecting the constraints on team composition, budget, and players per team. The overall algorithm for the FPL Belief State MDP is presented below:

3.5 Data Collection Methods

Implementing this project would not have been possible without clean, publicly-available data sources.

The most important data source was the gameweek-by-gameweek data in the data folder of my code repository that was cloned from the FPL Historical Dataset. The dataset is available in the Github repository [33].

While I later discovered that I could have retrieved season fixtures by manipulating the data from the forementioned repository, I ended up scraping fixture data from the unofficial Fantasy Premier League API using a Python script

I also scraped fixture results i.e. home team, away team, home team goals, and away team goals from football-data website

Algorithm 1 FPL Belief State MDP Algorithm

1: Initialize team $\mathbf{T}$ with players having highest points-per-cost
2: Initialize budget $B = B_0$
3: Initialize gameweek $GW = 1$
4: Initialize action subset with promising transfers
5: Initialize Bayesian Q-values $\mathcal{NG}(\mu_a, \lambda_a, \alpha_a, \beta_a)$ for each action
6: **while** $GW \leq 38$ **do**
7:     Select action $a = \arg\max_a\{\mu_a + VPI(a)\}$
8:     Execute transfer if specified by a
9:     Select captain with highest expected points
10:     Select bench players with lowest expected points while respecting formation rules
11:     Calculate reward (gameweek points - transfer cost)
12:     Update Bayesian Q-values using observed reward
13:     Update VPI for all actions
14:     Replace low-utility actions with new promising actions
15:     $GW = GW + 1$
16: **end while**

3.6 Technology Stack

3.6.1 Hardware Infrastructure

- Computational resources: Apple M2 Chip, 8 GB Ram, 256 GB Memory

- Computing environment: Local workstation

3.6.2 Software framework

- Operating System: macOS Sequoia 15.4

- Programming Languages: Python

- Integrated Development Environment (IDE): Visual Studio Code

- Version Control: Git, GitHub

3.6.3 Data Management

- Data Storage: GitHub, Local

- Data Format: CSV, Pickled Python objectives

- Data Processing Tools: Pandas

3.6.4 Analysis & Modeling

- Statistical Analysis Tools: Numpy, Pymc

- Machine Learning Libraries: Scipy, Gymnasium

- Visualization Tools: Matplotlib

- Domain-Specific Libraries: fpl [34]

3.6.5 Reproducibility Framework

- Environment Management: Miniconda virtual environment fpl_env

- Dependency Management: Miniconda

- Random Seed Control: Set $RANDOM\_SEED$ variable in constants.py

4 Results

Organize the result section according to major topics.

4.1 subsection

The reader will scan through the section headers at the first pass. Subsection headings help organizations and help the reader find the parts of interest.

Make section headings as specific and information-rich as possible. Make sure to interpret the results. Summarize the collection of experimental results as clearly and as economically as possible.

4.2 How you should prepare figures

Figure 4.1: QR Code for Poster Session Program, Spring 2024.

The figure and table captions should contain enough context so that a reader can understand the content of the figure or table without having to refer to the text.

Any labels or uncommon abbreviations need to be explained in the figure or table caption.

4.3 In the Results section..

- Organize this section according to major topics.

- Subheadings to make the organization clear and to help the reader

- scan through the text to find the parts of interest.

- Make these section headings as specific and information-rich as possible but try to cover the major common points.

- Interpret your results. (e.g., How do your results compare to related studies? Do your results accord with theory or are they surprising in some way? What are the underlying mechanisms that may explain what you found?)

- Summarize the collection of experimental results as clearly and as economically as possible.

- The figure and table captions should contain enough context so that a reader can understand the content of the figure or table without having to refer to the text.

- Any labels or uncommon abbreviations need to be explained in the figure or table caption.

5 Related Work

Avoid simply summarizing the findings of each paper, one paragraph per paper. [35] discusses..

Instead, synthesize the papers by identifying themes and use the papers to reference for the themes.

6 Conclusions

A good conclusion leaves readers with a clear statement of your point and a renewed appreciation of its significance.

Start with a brief statement with your research questions. Describe the main results. Provide a summary of the conclusion. Emphasize the contribution of your work.

A title is the first thing readers read and the last thing you should write. It is not just a few words to suggest the topic of your papers. A title is useful when it helps readers understand specifically what is to come. Put into your title the keywords in your main point.

7 Future Work

Bibliography

[1]   David Conn. *The Fall of the House of FIFA*. Random House, 2017.

[2]   Premier League. *Premier League Handbook 2023/24*. Tech. rep. Premier League, 2023.

[3]   Babatunde Buraimo and Rob Simmons. "Uncertainty of outcome or star quality? Television audience demand for English Premier League football". In: *International Journal of the Economics of Business* 22.3 (2015), pp. 449–469.

[4]   Tom Evens, Petros Iosifidis, and Paul Smith. "The political economy of television sports rights: History, power and culture". In: *International Journal of Sport Policy and Politics* 14.1 (2022), pp. 13–30.

[5]   Stefan Szymanski and Andrew Zimbalist. *National Pastime: How Americans Play Baseball and the Rest of the World Plays Soccer*. Brookings Institution Press, 2005.

[6]   Fantasy Premier League. *Official Fantasy Premier League Participation Statistics*. Tech. rep. Premier League, 2023.

[7]   Flavia Bonomo, Guillermo Durán, and Javier Marenco. "Mathematical programming as a tool for virtual soccer coaches: A case study of a fantasy sport game". In: *International Transactions in Operational Research* 21.3 (2014), pp. 399–414.

[8]   Fantasy Premier League. *Rules*. 2024. URL: https://fantasy.premierleague.com/help/rules (visited on 04/01/2024).

[9]   Mike Hughes and Ian Franks. *Analysis of Sport: The Key Concepts*. Routledge, 2005.

[10] Tom Decroos et al. "Actions speak louder than goals: Valuing player actions in soccer". In: *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. July 2019, pp. 1851–1861.

[11] Javier Fernández, Luke Bornn, and Dan Cervone. "A framework for the fine-grained evaluation of the instantaneous expected value of soccer possessions". In: *Machine Learning* 110.6 (2021), pp. 1389–1427.

[12] Quyen Tran and Seongsoo Lee. "Predicting fantasy premier league price changes using machine learning". In: *2022 International Conference on Information Networking (ICOIN)*. IEEE, 2022, pp. 622–627.

[13] Anthony C. Constantinou and Norman E. Fenton. "Towards smart-data: Improving predictive accuracy in long-term football team performance". In: *Knowledge-Based Systems* 124 (2017), pp. 93–104.

[14] Terence Matthews, Sarvapali D. Ramchurn, and Georgios Chalkiadakis. "Competing with humans at fantasy football: Team formation in large partially-observable domains". In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 26. 1. June 2012.

[15] Alina Bialkowski et al. "Large-scale analysis of soccer matches using spatiotemporal tracking data". In: *2014 IEEE International Conference on Data Mining*. IEEE, Nov. 2014, pp. 725–730.

[16] Alex Bryson, Bernd Frick, and Rob Simmons. "The returns to scarce talent: Footedness and player remuneration in European soccer". In: *Journal of Sports Economics* 14.6 (2013), pp. 606–628.

[17] Terence Matthews. "Improving fantasy football draft strategies using multiobjective optimization techniques". In: *Journal of Quantitative Analysis in Sports* 9.2 (2013), pp. 121–140.

[18] Anthony C. Constantinou. "Dolores: A model that predicts football match outcomes from all over the world". In: *Machine Learning* 108.1 (2019), pp. 49–75.

[19] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. 2nd ed. MIT Press, 2018.

[20] David Butler, Robert Butler, and John Eakins. "Expert performance and crowd wisdom: Evidence from fantasy premier league". In: *International Journal of Financial Studies* 9.1 (2021), p. 5.

[21] David Silver et al. "Mastering the game of go without human knowledge". In: *Nature* 550.7676 (2017), pp. 354–359.

[22] Terence Matthews and Sarvapali D. Ramchurn. "Competing with humans at fantasy football: Team formation in large partially-observable domains". In: *Journal of Autonomous Agents and Multi-Agent Systems* 33.2 (2019), pp. 130–171.

[23] Mark J. Dixon and Stuart G. Coles. "Modelling association football scores and inefficiencies in the football betting market". In: *Journal of the Royal Statistical Society: Series C (Applied Statistics)* 46.2 (1997), pp. 265–280.

[24] Giovanni Pantuso. "The football team composition problem: A stochastic programming approach". In: *Journal of Quantitative Analysis in Sports* 13.3 (2017), pp. 113–129.

[25] Alexander P. Rotshtein, Morton Posner, and Anna B. Rakityanskaya. "Football predictions based on a fuzzy model with genetic and neural tuning". In: *Cybernetics and Systems Analysis* 41.4 (2005), pp. 619–630.

[26] Chuan Guo et al. "On calibration of modern neural networks". In: *International Conference on Machine Learning*. PMLR, July 2017, pp. 1321–1330.

[27] Rahul Baboota and Harleen Kaur. "Predictive analysis and modelling football results using machine learning approach for English Premier League". In: *International Journal of Forecasting* 35.2 (2019), pp. 741–755.

[28] Ondřej Hubáček, Gustav Šourek, and Filip Železný. "Deep learning from label proportions for sequential data". In: *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*. Cham: Springer, May 2019, pp. 39–55.

[29] Mohammad Amin Rahimian et al. "Minimizing maximal regret in commitment games". In: *AAAI Conference on Artificial Intelligence*. Vol. 35. 6. 2021, pp. 5493–5500.

[30] Luca Pappalardo et al. "A public data set of spatio-temporal match events in soccer competitions". In: *Scientific Data* 6.1 (2019), pp. 1–15.

[31] Tom Decroos and Jesse Davis. "Player vectors: Characterizing soccer players' playing style from match event streams". In: *Machine Learning and Knowledge Discovery in Databases: European Conference, ECML PKDD 2019*. Springer International Publishing, 2020, pp. 569–584.

[32] David Sheehan. *Predicting Football Results With Statistical Modelling: Dixon-Coles and Time-Weighting*. https://dashee87.github.io/football/python/predicting-football-results-with-statistical-modelling-dixon-coles-and-time-weighting/. 2018.

[33] Vaastav Anand. *FPL Historical Dataset*. Retrieved August 2022 from https://github.com/vaastav/Fantasy-Premier-League/. 2022.

[34] David MacLeod and Hayden MacDonald. *A Python wrapper for the Fantasy Premier League API*. https://fpl.readthedocs.io/en/latest/index.html. 2019.

[35] Mohammad Hosseini and Viswanathan Swaminathan. "Adaptive 360 VR video streaming: Divide and conquer". In: *2016 IEEE International Symposium on Multimedia (ISM)*. IEEE. 2016, pp. 107–110.

A   Appendix Title

Instructions on how to reproduce your experiment. Anything that is not covered at the main but you want to add. Additional plots/results, etc.